

## **Deep Learning – GAN Assignment**

### **Part 1**

The discriminator's input is of the data dimensions (e.g., 8 for diabetes without the label) and its architecture contains 3 dense layers – The first layer with 256 units and an activation function of LeakyReLU, the second layer containing 512 units and an activation function of LeakyReLU and the last layer contains 1 unit with an activation function of sigmoid. After the first layer, we added a dropout of 0.2. This way, the discriminator informs us about whether the sample is real or fake – 1 if real and otherwise 0.

The generator's input is of latent dimensions (of the noise it will get, e.g., 8 in Diabetes data and 20 in German Credit data) and its architecture contains 3 dense layers – The first layer with 512 units and activation function of LeakyReLU, the second layer containing 256 units with an activation function of LeakyReLU and the last layer contains number of units as the number of the data dimensions(e.g. 8 for Diabetes data) with an activation function of tanh. After the first layer, we added a dropout of 0.4.

This architecture and the chosen hyperparameters are based on this paper <https://arxiv.org/pdf/1904.09135.pdf> which used on the Diabetes data.

Our GAN model's outputs are in the scale of -1 to 1 (the last layer's activation function is tanh), thus for the diabetes data, we first performed min-max scaling in order to be able to produce legitimate examples. For the German credit data, for the categorical features, we created one hot vectors and for the ordinal features we used ordinal encoder. Then in order to scale all the values between -1 to 1 we used min max scaler with appropriate range. After the fake samples' generation, we inverse transform the values back to their classes.

We created 2 GAN models – one for the diabetes data and one for the German credit data.

## Questions answering - Analyzing the products of our model:

**3.A.** We calculated the distance between the original samples and the generated ones and chose the closest ones after transformation. Then we used inverse transformation to reconstruct the samples. (We also cast the samples into integers for better understanding).

### Diabetes Data

#### **Samples that fooled the detector:**

```
sample number 24:
Fake sample : [ 3 115 77 38 212 33 0 29]
Closest real : [ 2 104 80 45 191 33 0 29]
Euclidean Distance: 0.3078351318836212
sample number 28:
Fake sample : [ 2 152 81 54 318 41 0 24]
Closest real : [ 1 136 74 50 204 37 0 24]
Euclidean Distance: 0.4348710775375366
sample number 10:
Fake sample : [ 1 110 78 39 91 34 0 22]
Closest real : [ 2 107 74 30 100 33 0 23]
Euclidean Distance: 0.2132774293422699
sample number 76:
Fake sample : [ 0 171 81 32 108 41 0 21]
Closest real : [ 0 162 76 36 0 49 0 26]
Euclidean Distance: 0.47395753860473633
sample number 12:
Fake sample : [ 1 80 77 16 15 23 0 26]
Closest real : [ 1 81 72 18 40 26 0 24]
Euclidean Distance: 0.22155877947807312
```

As we can see the values that fooled the detector are very similar to the real ones. Also, the distances between fake examples and real ones are low.

#### **Samples that did not fooled the detector:**

```
sample number 52:
Fake sample : [ 13 150 59 41 20 36 0 27]
Closest real : [ 14 175 62 30 0 33 0 38]
Euclidean Distance: 0.5134834051132202
sample number 82:
Fake sample : [ 9 145 62 25 43 29 0 31]
Closest real : [ 10 125 70 26 115 31 0 41]
Euclidean Distance: 0.4571422338485718
sample number 95:
Fake sample : [ 15 143 44 15 0 32 0 25]
Closest real : [ 14 175 62 30 0 33 0 38]
Euclidean Distance: 0.689517080783844
sample number 75:
Fake sample : [ 11 124 58 31 21 32 0 33]
Closest real : [ 10 129 62 36 0 41 0 38]
Euclidean Distance: 0.3906545341014862
sample number 27:
Fake sample : [ 9 158 61 42 73 36 0 27]
Closest real : [ 10 129 62 36 0 41 0 38]
Euclidean Distance: 0.5548930168151855
```

As we can see the values that did not fool the detector are a little bit different from the real ones. Also, the distances between fake examples and real ones are higher than the previous samples that fooled the detector.


In addition, we calculated the mean distance between fake samples which fooled the detector and their closest real samples, and the mean distance between fake samples which **did not** fool the detector and their closest real samples. We found out that the first mean distance ("fooled") is lower than the second one (Did not "fool"):

The mean distance of samples which fooled the detector: 0.3738189935684204


---

The mean distance of samples which did not fool the detector: 0.47054868936538696

The following visualization shows the real samples (red), the samples that fooled the detector (blue) and the samples that did not fool the detector (black). As we can see the samples that fooled the detector (blue) are indeed similar to some of the original data. But the samples that did not fool the detector are very close to samples that did fool, so it is very hard to say why the model decided to classify these samples as fake.




**3.B.** We generated 100 samples at random after our model was converged. The result showed 82 of 100 generated samples were able to pass as real samples. In the following graphs we used a PCA dimensionality reduction technique to visualize the real and fake data. As we can see, our model was able to generate samples very similar to the real data. The red dots represent real data, the blue ones represent fake data (1000 samples), and the yellow ones represent the closest real and fake samples. The left graph shows the real and fake data, while the right one shows only the real data.



The generated samples are presented on the notebook.

**3.C.** Graph describing the loss of the generator and the discriminator:



As we can see from the graph, in the beginning of the training both models go "back and forth" while the discriminator was far better than the generator (the loss of the generator is very high).

The discriminator was a consistent leader while its loss is between 0.4-0.6 and the loss of the generator is a little bit higher between 0.7-1.

We can understand the result because the classification task of the discriminator is easier than the generator task.

## German Credit Data

### Samples that fooled the detector:

```
sample number 29:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 2 0 0 0
3 14 1697 1 3 24 1 1]
Closest real : [ 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
0 0 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0
1 1 0 0 0 0 0 1 0 1 0 1 0 3 2 0 0 0
2 6 1236 2 4 50 1 1]
Euclidean Distance: 3.1874539852142334
sample number 38:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 2 0 0 0
3 13 1299 3 3 49 3 1]
Closest real : [ 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0
0 0 0 1 0 1 0 0 1 0 0 0 0 0 0 0 0 0
1 0 1 0 0 1 0 0 1 0 1 0 3 0 0 0 0 0
4 11 1254 4 4 61 2 1]
Euclidean Distance: 2.2955758571624756
sample number 32:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 3
2 6 1759 2 2 28 1 1]
Closest real : [ 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0
0 0 1 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0
1 0 1 0 0 1 0 0 1 0 1 0 1 0 4 0 0 0
2 8 3194 1 2 33 1 1]
Euclidean Distance: 3.2007484436035156
sample number 85:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 3 13 696 3 2 27 1 1]
Closest real : [ 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 1
0 1 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0
1 0 1 0 0 0 4 11 651 4 4 24 1 1]
Euclidean Distance: 3.593585729598999
```

As we can see the values that fooled the detector are very similar to the real ones. Also, the distances between fake examples and real ones are bigger than the previous distances in diabetes data, and that's because we have more features here and thus we have more arguments in the distance function what it makes it bigger.

### Samples that did not fool the detector:

```
sample number 46:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0
3 9 1336 3 3 47 1 1]
Closest real : [ 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0
0 0 0 1 0 1 0 0 0 0 0 0 1 0 0 0 0 0
1 0 0 1 0 0 1 0 1 0 1 0 1 0 1 1 0 0
4 23 3068 4 4 30 1 1]
Euclidean Distance: 2.3567824363708496
sample number 19:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 2 3 3 11 699 3 3 26 1 1]
Closest real : [ 0 0 1 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0
0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0
1 0 1 0 0 0 1 0 1 0 1 0 1 0 4 0 0 0
2 18 1042 4 2 33 1 1]
Euclidean Distance: 3.29418683052063
sample number 17:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 2 0 0 0
1 9 1189 3 3 25 3 1]
Closest real : [ 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0
0 0 1 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0
1 0 1 0 0 0 1 0 0 1 1 0 3 0 0 0 0 0
2 30 2831 4 2 33 1 1]
Euclidean Distance: 3.6620266437530518
sample number 91:
Fake sample : [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 1 0 3 14 655 3 3 39 2 1]
Closest real : [ 0 0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0
0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 0 0
0 0 1 0 0 0 1 0 1 0 1 0 1 0 2 0 0 0
2 18 2100 4 2 37 1 1]
Euclidean Distance: 2.9420056343078613
```

As we can see the values that did not fool the detector are different from the real ones.

In addition, we calculated the mean distance between fake samples which fooled the detector and their closest real samples, and the mean distance between fake samples which **did not** fool the detector and their closest real samples. We found out that both distances are very close


The mean distance of samples which fooled the detector: 3.4326775074005127

---


The mean distance of samples which did not fool the detector: 3.432727098464966

---

The following visualization shows the real samples (red), the samples that fooled the detector (blue) and the samples that did not fool the detector (black). As we can see the samples that fooled the detector (blue) are indeed similar to some of the original data. But the samples that did not fool the detector are very close to samples that did fool and to real samples, so it is very hard to say why the model decided to classify these samples as fake.




**3.B.** We generated 100 samples at random after our model was converged. The result showed 29 of 100 generated samples were able to pass as real samples. In the following graphs we used a PCA dimensionality reduction technique to visualize the real and fake data. It seems that our model was able to generate samples very similar to the real data, but the model was able to detect them. The red dots represent real data, the blue ones represent fake data (1000 samples). The left graph shows the real and fake data, while the right one shows only the real data.



The generated samples are presented on the notebook.

**3.C.** Graph describing the loss of the generator and the discriminator:



As we can see from the graph, in the beginning of the training both models go "back and forth" while the discriminator was far better than the generator (the loss of the generator is higher).

The discriminator was a consistent leader while its loss is between 0.4-0.7 and the loss of the generator is a little bit higher between 0.6-1.

We can understand the result because the classification task of the discriminator is easier than the generator task.